

Research summary

Maria Paola Bonacina

(March 29, 2025)

My research area is *automated reasoning*, or how to make computers reason, not necessarily like humans, but rather in their own way. Reasoning problems involve *validity* queries (does a *conjecture* φ follow from a set H of *assumptions*?) or *satisfiability* queries (does a set S of constraints admit solution?), where formulæ express properties of an object of study (e.g., a system, a program, a datatype, a circuit, a protocol, a mathematical structure). The answer to a validity query is a *proof* (that φ follows from H or that $S = H \cup \{\neg\varphi\}$ is unsatisfiable) or a *counter-example* (a *model* of S). The answer to a satisfiability query is a *model* (of S) or a *proof* (that S is unsatisfiable). Automated reasoning is about designing methods to solve these problems, proving properties of such methods (e.g., *soundness*, *completeness*, *termination*), and implementing them in *automated reasoners*, interfaced with human or software users.

My research is motivated by both fundamental challenges and applications. A constant challenge is improving the trade-off between *generality* (how large is the class of problems a method can handle) and *efficiency*. The use of automated reasoning for the *analysis*, *verification*, *synthesis*, and *optimization of software* is a natural application, because logic is the calculus of computation [46]. Automated reasoners are used for discharging verification or synthesis conditions, refining abstractions, generating tests for testing, and examples for synthesis. Correct-by-construction software, provable privacy, verification of distributed protocols and systems, and verification of randomized algorithms are among the current challenges in this area. Reasoners play an increasingly crucial rôle in ensuring features, such as reliability and privacy, that are of extreme importance for society.

The application of automated reasoning to *improve generative AI techniques* is another challenge [41]. An even more fundamental challenge is the *combination of automated reasoning and generative AI techniques* towards a more sophisticated artificial intelligence. Automated reasoning is complementary to machine learning, as automated reasoning allows the machine to reason based on general laws or principles, whereas machine learning allows the machine to learn from data. Automated reasoning is connected with symbolic computation (e.g., constraint problem solving, computer algebra), and computational logic (e.g., declarative programming, rewriting), and it finds applications also in deductive databases, declarative programming, mathematics, and education.

My research to date can be presented in four overlapping threads:

- A. Theorem-Proving Strategies and Satisfiability Procedures**
- B. Interpolation of Proofs**
- C. Distributed Automated Deduction**
- D. Strategy Analysis**

summarized in the sequel, with citations referring to the publications in my curriculum vitae.

A. Theorem-Proving Strategies and Satisfiability Procedures

Most reasoning methods transform the problem $H \cup \{\neg\varphi\}$ into an equisatisfiable set S of *clauses*, a standard machine format. In first-order logic (FOL) unsatisfiability is semidecidable, satisfiability is not even semidecidable, and reasoning methods are *semidecision procedures* called *theorem-proving strategies* [91, 76, 87, 74, 73, 5, 4] and implemented in *theorem provers*. A theorem-proving strategy is characterized by an *inference system* and a *search plan*. An inference system is a set of *inference rules* that manipulate clauses, until a *refutation* is found, and a *proof* can be reconstructed from the *derivation*. A strategy may reason mostly *forward* (i.e., from the assumptions H) or mostly *backward* (i.e., from the clauses in the clausal form of $\neg\varphi$, called *goal clauses*), to the point of being *goal-sensitive*, if all generated clauses are connected with a goal clause. A strategy may be *semantically-guided* by a given *fixed* interpretation, and it is *proof confluent*, if it does not need to undo inferences by backtracking.

Ordering-based strategies work with a set of clauses, initially the input set S . *Expansion* inference rules, such as *resolution*, *paramodulation*, and *superposition*, generate and add clauses, consequences of the existing ones. *Contraction* inference rules, such as *subsumption* and *simplification*, delete or replace *redundant* clauses. A refutation is reached when the empty clause \square is generated. *Well-founded orderings* on terms, literals, clauses, or proofs, are used to restrict expansion, and to define contraction and redundancy. Contraction inference rules and an *eager-contraction* search plan characterize *contraction-based* strategies, that are a default choice when *equality* is involved. Typical ordering-based strategies reason primarily *forward*, as $\neg\varphi$ is treated as an additional hypothesis, may be *semantically-guided*, and are *proof confluent* [91, 76, 73, 4].

Subgoal-reduction strategies, based on *linear resolution*, *model elimination*, or *tableaux*, apply inferences to reduce a current goal to subgoals [91, 76, 74, 73]. They operate on a stack of goals and use *depth-first search with backtracking and iterative deepening*. The stack of goals is the *frontier* of a tree-like structure, a tableau, where branches represent possible models. A refutation is found when all branches are *closed* as contradictory. *Lemmaizing* (i.e., turning solved goals into lemmas) and *caching* (i.e., storing solved or failed goals in a look-up table) counter the redundancy of repeated subgoals. Typical subgoal-reduction strategies reason mostly *backward* and are *goal-sensitive*. *Instance-based strategies* generate *instances* of clauses, and invoke a satisfiability procedure to test sets of ground instances for unsatisfiability. Typical instance-based strategies reason mostly *forward*, and are *model-driven*, if they generate instances that are false in the model found by the satisfiability procedure when it detects satisfiability.

In propositional logic and in decidable fragments of FOL or of a first-order theory \mathcal{T} , satisfiability is decidable, and reasoning methods are *decision procedures* called *satisfiability procedures* and implemented in *satisfiability solvers*. A satisfiability procedure is characterized by a *transition system* and a search plan. A transition system is a set of *transition rules* that transform a *trail* representing a *candidate model*, until either a model is found or an unsolvable *conflict* reveals that no model exists. These procedures are *model-based*, as they build and discard candidate models, and *conflict-driven*, as they apply nontrivial inferences only to *explain* and *solve* conflicts [39, 86].

The archetypal satisfiability procedure is the *Conflict-Driven Clause Learning* (CDCL) procedure for propositional logic. It works by *deciding* truth assignments to atoms and *propagating* their

consequences. When a conflict arises (a clause is false in the current assignment), the procedure *explains* it by *resolution* and *learns* a *lemma* to avoid hitting that conflict again. The $\text{CDCL}(\mathcal{T})$ procedure integrates a satisfiability procedure for a theory \mathcal{T} in CDCL. If \mathcal{T} is a union of theories $\mathcal{T}_1, \dots, \mathcal{T}_n$, the existence of a satisfying \mathcal{T} -model depends on whether there are \mathcal{T}_i -models that agree on the interpretation of shared symbols and on the cardinalities of shared sorts. The *Nelson-Oppen scheme*, originally named *equality sharing*, assumes that the theories are *stably infinite*, meaning that they admit models with countably infinite domains for all sorts, and *disjoint*, meaning that they do not share function or predicate symbols other than equality, so that the \mathcal{T}_i -models only need to agree on which shared constants (or free variables) are equal. The equality sharing method combines the \mathcal{T}_i -satisfiability procedures as *black-boxes* that only need to propagate all entailed disjunctions of equalities between shared constants [72]. CDCL was generalized to *conflict-driven satisfiability procedures* for quantifier-free fragments of arithmetic. Key features of such procedures are *assignments to first-order variables* and conflict explanation by lemmas that may contain *new* (i.e., non-input) atoms. The MCSAT procedure integrates CDCL and one conflict-driven theory satisfiability procedure [74, 39, 72].

My research on theorem-proving strategies and satisfiability procedures develops *cross-fertilization methods* that unite features from different paradigms, while understanding inference rules and transition rules as *transformation of candidate proofs or candidate models*.

A1. Target-Oriented Completion. I investigated how to make ordering-based strategies *goal-sensitive*, or equivalently *target-oriented* (φ is the *target* theorem) in the context of *completion procedures* [110, 71, 101, 68, 109]. Completion was understood as generation of confluent rewrite systems for equational theories. Theorem proving was either a side-effect or the successive application of the confluent rewrite system to decide the word problem by rewriting. The latter way is impractical, since few equational theories have finite confluent rewrite systems. I understood that the first one is intrinsically inefficient, because in order to generate a confluent system, or a *saturated set* in FOL, the procedure performs inferences that are *unnecessary* to prove φ . The key point is *fairness*. The pre-existing notion, *uniform fairness*, captures the inferences needed to generate a saturated set. I proposed a new notion of *fairness* that captures the inferences needed to prove φ . My *target-oriented* framework for completion is based on applying *proof orderings* to the proofs of φ [68, 109, 67, 100, 108, 26, 24]. It covered all known completion procedures for equational logic, including practical *target-oriented techniques* [106, 69, 65], *inductionless induction*, and the generation of confluent rewrite systems as a special case. In experiments, I obtained the *first automated proof of the “Dependency of the Fifth Axiom” in Łukasiewicz’s many-valued logic* [107, 69, 99], introducing it as a source of challenge problems. The *Linear Completion* procedure that interprets *rewrite programs* is another instance of target-oriented completion. I defined the *operational* and *denotational semantics* of rewrite programs, disproving the folklore that they are the same as Prolog, showing the different expressive power of programming with *bi-implications* versus *implications*, and the effect of *simplification* on *termination* [70, 109, 30]. I also gave *counter-examples to the completeness of the RUE/NRF inference systems* [98].

A2. Lemmatization from Model Elimination to Semantic Resolution. Understanding that *lemmaizing* is a form of forward reasoning in subgoal-reduction strategies, I had the idea of using it to increase the *goal-sensitivity* of *semantic resolution* strategies [59, 21], where semantic guidance orients the strategy towards forward or backward reasoning. Lemmaizing is a *meta-level inference rule*, because it derives a lemma based on a whole fragment of the derivation. I showed that it can be used to add backward inferences to a forward strategy and vice versa. I defined a set of inference rules that implement lemmaizing in semantic resolution strategies, and I showed how to add contraction, including *purity deletion*, to such strategies. Thus, lemmaizing and contraction can coexist, and contraction can take advantage of unit lemmas. For subgoal-reduction strategies, I formalized *caching* and *depth-dependent caching* as inference rules justified by the meta-rules for lemmaizing. I observed that lemmaizing and caching allow subgoal-reduction strategies to keep some generated clauses, and therefore make *subsumption* possible, bringing a feature of ordering-based strategies to subgoal reduction.

A3. Inferences and Canonicity. I expanded my early work on target-oriented completion in an analysis of how inferences *reduce proofs* and *transform presentations* [16]. A presentation is *contracted*, if it is made of the premises of the *minimal proofs*; *canonical*, if it is made of the premises of the minimal proofs in the whole theory (*normal form proofs*); *complete*, if it offers at least a normal form proof for each theorem; and *saturated*, if it features all normal form proofs for all theorems. Therefore, canonical and saturated coincide only if normal form proofs are unique, and *complete*, rather than *saturated*, is sufficient for theorem proving. Accordingly, while a *uniformly fair* derivation produces a *saturated* presentation, a *fair* derivation only yields a *complete* one. In practice, a search plan should schedule enough expansion and contraction inferences to get in the limit a complete and contracted presentation. I applied this framework to *implicational systems*, that are presentations of *propositional Horn theories* [48, 75]. Given an implicational system \mathcal{T} and a set X of propositional variables, the problem is to find the least \mathcal{T} -model that satisfies X . Implicational systems can be translated into propositional rewrite systems, where the rewrite relation is bi-implication as in Linear Completion. I defined a completion procedure for propositional rewrite systems that computes the least \mathcal{T} -model of X and transforms \mathcal{T} into a *canonical* equivalent presentation. I also studied canonicity in *conditional equational theories*, where proof normalization yields decision procedures based on saturated presentations [75].

A4. Composing Theorem-Proving Inference Systems and Satisfiability Procedures

If a theorem-proving strategy is *guaranteed to halt* on problems of a certain kind, it is a *decision procedure* for that class of problem. Applying theorem-proving strategies to *satisfiability modulo theories* (SMT) problems offers several benefits: existing theorem provers can be used *off the shelf*, correctness and completeness are given, combination of theories reduces to giving as input the union of their presentations, and theorem provers generate proofs. With this motivation, I produced a number of results showing that a standard *ordering-based inference system*, known in the literature as the *superposition calculus*, or even *superposition* for short, generates only finitely many clauses from certain satisfiability problems, so that every fair strategy with that inference system is a decision procedure for those problems [82, 81, 80, 53, 52, 104, 14, 51, 50, 103, 90, 15].

On the other hand, theories such as arithmetic or bitvectors, do not lend themselves to reasoning by generic inferences. Thus, I investigated first how to *pipeline* [49, 13] and then how to *integrate* [47, 12] the superposition calculus and the CDCL(T) procedure.

A4.1. Superposition Decision Procedures. We showed that superposition decides the satisfiability of sets of ground literals in the theories of *records with or without extensionality, possibly empty lists, arrays with or without extensionality, integer offsets, integer offsets modulo* [80, 53, 14], and *recursive data structures* [51], including *acyclic non-empty lists* as a special case. I discovered a condition, called *variable-inactivity*, whereby if the theories are *disjoint* and *variable-inactive*, and superposition terminates on satisfiability problems in each theory, then it terminates also on satisfiability problems in their union [53, 14]. All the above mentioned theories are variable-inactive. Contrary to the folklore that a generic prover cannot compete with solvers with built-in theories, the experimental comparison of the *E prover* with the CVC and CVC Lite SMT-solvers was overall favorable to the theorem prover [80, 81, 53, 14]. We also showed that if a theory is *not stably infinite*, superposition is guaranteed to generate eventually an *at-most cardinality constraint*, so that the theory is *not variable-inactive* [52, 104]. Thus, *variable-inactivity implies stable-infiniteness*, and superposition can discover the *lack of infinite models* by generating an at-most cardinality constraint [52, 14]. Our superposition decision procedures for *records without extensionality* and *integer offsets modulo* [51, 14], and for *integer offsets* and *records with extensionality* [90, 15], are *polynomial*, and the one for the latter theory was the first with this property. Next, we showed that for *variable-inactive* theories, if superposition decides the satisfiability of sets of ground literals, it also decides the satisfiability of sets of ground clauses [90, 15]. This result applies to the theories of *equality, non-empty possibly cyclic lists, arrays with or without extensionality, injective arrays* [50, 103], *finite sets with or without extensionality, records with or without extensionality, possibly empty possibly cyclic lists, integer offsets modulo, recursive data structures*, and all their unions.

A4.2. Pipelining Superposition and CDCL(T). In problems from applications the input clause set S may contain very long clauses, and while CDCL-based solvers break clauses apart by case analysis, superposition generates longer and longer clauses. *Decision procedures by stages* addresses this obstacle by pipelining superposition with an SMT-solver [49, 13]. S is partitioned into a set S_1 of unit clauses and a set S_2 of non-unit clauses. Superposition saturates $\mathcal{T} \cup S_1$ into $\mathcal{T} \cup \bar{S}$, where \bar{S} is finite, ground, and capable of entailing all clauses that can be generated from $\mathcal{T} \cup S_1$, and $\bar{S} \cup S_2$ is fed to the SMT-solver. We found sufficient conditions to ensure that \bar{S} has these properties, and we obtained *\mathcal{T} -decision procedures by stages* for *arrays with or without extensionality, records with or without extensionality, integer offsets*, and their unions. Here saturation works as an inference-based *reduction to the theory of equality*. If the problem involves two theories \mathcal{T}_1 and \mathcal{T}_2 , such that superposition is a decision procedure for each, we decompose S into $\mathcal{T}_1 \cup S_1$, $\mathcal{T}_2 \cup S_2$ and S_3 , where S_1 contains unit \mathcal{T}_1 -clauses, S_2 contains unit \mathcal{T}_2 -clauses, and S_3 contains the remaining clauses. The procedure saturates $\mathcal{T}_1 \cup S_1$ into $\mathcal{T}_1 \cup \bar{S}_1$ and $\mathcal{T}_2 \cup S_2$ into $\mathcal{T}_2 \cup \bar{S}_2$, and passes $\bar{S}_1 \cup \bar{S}_2 \cup S_3$ on to the SMT-solver. Thus, the part of the problem involving theories such as arithmetic or bitvectors can be given directly to the SMT-solver.

A4.3. The CDCL($\Gamma + \mathcal{T}$) Procedure with Speculative Inferences. In problems from applications the input clause set S often contains ground clauses with \mathcal{T} -symbols and a subset \mathcal{R} of non-ground clauses without \mathcal{T} -symbols. The CDCL($\Gamma + \mathcal{T}$) procedure integrates a superposition-based inference system Γ in CDCL(\mathcal{T}), in such a way that Γ works with non-ground \mathcal{R} -clauses and ground unit \mathcal{R} -clauses on the trail, while CDCL(\mathcal{T}) takes care of ground clauses [47, 12, 44]. Since trail literals may be withdrawn upon backjumping, they are memorized in clauses as *hypotheses*, that are inherited through inferences. When backjumping removes literals from the trail, the clauses depending on them are also removed. Contraction rules are adjusted to take this dynamic effect into account. If \mathcal{R} is *variable-inactive*, CDCL($\Gamma + \mathcal{T}$) is *refutationally complete* [47, 12]: indeed, since variable-inactivity implies stable infiniteness [52], and superposition is guaranteed to generate clauses that entails all disjunctions of equalities between shared constants [13], the completeness requirements for equality sharing are fulfilled. As CDCL(\mathcal{T}) uses depth-first search with backtracking, the *fairness* of CDCL($\Gamma + \mathcal{T}$) requires *iterative deepening* on the number of Γ -inferences. If S is unsatisfiable, CDCL($\Gamma + \mathcal{T}$) is guaranteed to halt with a contradiction; otherwise, it may either halt with a model, or get *stuck* at the current limit on number of Γ -inferences. The third outcome is excluded for those theories for which Γ is a decision procedure. In order to get more decision procedures, CDCL($\Gamma + \mathcal{T}$) features *speculative inferences*: it can add to the current set an arbitrary clause, with as hypothesis a new propositional variable added to the trail to keep track of the decision [47, 12]. If S is satisfiable, and the added clause causes a contradiction, CDCL($\Gamma + \mathcal{T}$) handles it as a conflict, undoing the speculative inference by backjumping. If we can provide a sequence of clauses (e.g., equalities) whose addition enforces termination, CDCL($\Gamma + \mathcal{T}$) is a decision procedure. This is the case for several *axiomatizations of type systems* [47, 12].

A5. SGGS: a Satisfiability Procedure for FOL

We designed SGGS to be the first theorem-proving method that is *simultaneously* first-order, *semantically guided*, *goal-sensitive*, *model-based*, *conflict-driven*, and *proof confluent* [88, 78, 43, 9, 8]. SGGS is the first method that succeeded in *generalizing CDCL to FOL*.

A5.1. Model Representation in SGGS. SGGS is *semantically guided* because it assumes a fixed *initial interpretation* I . Given I and input set S of clauses, if $I \models S$, the problem is solved. Otherwise, SGGS seeks to build a model of S by determining which literals that are true in I , called *I -true* literals, should be falsified to satisfy S . The current candidate model is represented by a trail Γ , which is a sequence of (possibly constrained) *non-ground* clauses with *selected literals* [9]. A key idea in SGGS is that *all ground instances of a literal march in lockstep*, and this is the reason for introducing constraints (e.g., if all ground instances of $P(x)$ are true except $P(b)$ we can write $x \neq b \triangleright P(x)$) [78]. Since the variables in first-order clauses, and hence literals, are universally quantified, if L is true all its ground instances are, but it suffices that one ground instance is false to make L false. We say that L is *uniformly false* if all its ground instances are, that is, if its flip $\neg L$ is true. We call *I -false* a literal that is uniformly false in I . SGGS builds Γ in such a way that all literals in all clauses in Γ are either *I -true* or *I -false*, and *I -false* literals are preferred for selection. An *I -true* literal is selected only in an *I -all-true* clause, that is, a clause such that all its literals are *I -true*. The associated interpretation $I[\Gamma]$ is I modified to satisfy the

selected literals in Γ . Thus, *literal selection* plays the role of *decision* in CDCL. For *first-order clausal propagation*, a literal L is uniformly false in $I[\Gamma]$, if all its ground instances appear negated among those that a selected literal M makes true in $I[\Gamma]$. If L is I -true, SGGS *assigns* it to (the clause of) M . A clause C is a *conflict clause* if all its literals are uniformly false in $I[\Gamma]$. If all literals in C , except the selected literal L , are uniformly false in $I[\Gamma]$, literal L is *implied* and C is its *justification*. SGGS ensures that every I -all-true clause in Γ is either a *conflict clause* (all its literals are assigned) or the *justification* of its selected literal (all its literals are assigned except the selected one). Since these assignments are computed in the SGGS inferences, there is no need of a first-order two-watched literal scheme to perform first-order clausal propagation [9].

A5.2. The SGGS Transition System. In an *SGGS-derivation* each trail is generated from its predecessor and S by applying an SGGS rule [8]. *SGGS-extension* adds to the trail an instance of an input clause and selects one of its literals. The added instance is built in order to capture ground instances of the input clause not satisfied by the current $I[\Gamma]$. *SGGS-deletion* deletes clause C from $\Gamma \text{CT}'$, if C is satisfied by $I[\Gamma]$. Similar to CDCL, if SGGS-extension adds to Γ a conflict clause E , *SGGS-resolution* *explains* the conflict resolving upon an I -false literal in E and the I -true selected literal of a justification. SGGS ensures that all I -false literals in E can be resolved away in this manner, yielding either \square or an I -all-true conflict clause C . *SGGS-move* solves the conflict and *learns* C , by moving it to the left of the clause whose selected literal makes C 's selected literal uniformly false: C 's selected literal becomes an implied literal. Thus, SGGS gets out of conflict without undoing inferences by backtracking or backjumping. *SGGS-splitting* of clause C by clause D replaces C by a *partition*, where all ground instances that a specified literal in C has in common with D 's selected literal are confined to one element. This enables SGGS-resolution or SGGS-deletion to remove such *intersections*, ridding Γ of contradictions or duplications. SGGS *makes progress* in two ways: either it extends Γ by an SGGS-extension, or it repairs $I[\Gamma]$ by either explaining and solving a conflict or removing an intersection in Γ . *Fairness* ensures that SGGS-deletion and other clause removals are applied eagerly, trivial splitting is avoided, progress is made whenever possible, every SGGS-extension generating a conflict clause is *bundled* with explanation and conflict-solving inferences to solve the conflict before further extensions, and inferences applying to shorter prefixes of the trail are never neglected in favor of others applying to longer prefixes. SGGS is *refutationally complete* and *model complete* in the limit (if the input is satisfiable, the limit of every fair SGGS-derivation represents a model) [8].

A5.3. SGGS, Decision Procedures, and Horn Theories. By model completeness in the limit *SGGS decision procedures are model-constructing*. We proved that SGGS decides several known decidable fragments of FOL: *stratified* [37, 3], *positive variable dominated* (PVD) [37, 3], *bounded depth increase* (BDI) [3], and *Datalog* [3]. The stratified fragment is a many-sorted generalization of the *Bernays-Schönfinkel class*, whose clausal version is known as *Effectively Propositional logic* (EPR). On the other hand, SGGS with *sign-based semantic guidance* (i.e., I is either *all-negative* – all negative literals are true – or *all-positive* – all positive literals are true) does *not* decide other known decidable fragments of FOL: *Ackermann*, *monadic*, FO^2 , and *guarded* [3]. These counterexamples show that the existence of a finite model does *not* imply

the termination of SGGS with sign-based semantic guidance. Other examples show that SGGS terminates and represents with a finite trail an infinite Herbrand model, so that termination does *not* imply the existence of a finite Herbrand model. We discovered several new decidable fragments of FOL by showing that SGGS decides them: *positively/negatively restrained*, *positively/negatively sort-restrained*, and *sort-refined-PVD*. Furthermore, since the size of SGGS-generated models can be upper-bounded, these new fragments enjoy the *small model property*. Since restrainedness is an ordering-based property, it can be reduced to termination of rewriting: this means that it is undecidable in general, but in practice termination tools can be applied to find restrained sets [37, 3]. We also investigated the behavior of SGGS on Horn clauses [34], showing that SGGS with all-negative I generates the least fixpoint model of a set of definite clauses, and the first negative conflict clause announces a refutation. Given Horn clauses, SGGS with all-negative I reasons forward, and SGGS with all-positive I reasons backward. The SGGS prototype Koala exhibited promising experimental results, especially on satisfiable problems [37, 34, 3].

A6. CDSAT: Conflict-Driven SATisfiability Modulo Theories and Assignments

CDSAT is a *conflict-driven method for reasoning in theory unions* that generalizes CDCL, equality sharing, CDCL(T), and MCSAT, solving the problem of combining multiple theory procedures, whether conflict-driven or not [42, 40, 77, 38, 7, 36, 6, 35, 1].

The CDSAT Framework. CDSAT treats formulæ as terms of sort **prop** (for proposition), and works with *assignments* of values to terms, including both Boolean and first-order terms, hence *Boolean assignments* and *first-order assignments*. Assignable values (e.g., truth values, numbers) are represented by special constants, introduced by *conservative theory extensions*, so that terms and values remain separate [42, 7]. Input problems are also assignments: an SMT problem is given by a Boolean assignment, while an input problem containing also first-order assignments is an SMA problem. SMA stands for *Satisfiability Modulo theories and Assignments*. An SMA problem asks whether there is a model of the input formula that incorporates the input first-order assignments. SMA problems arise in the reduction of an optimization problem to iterative satisfiability, where the input first-order assignments exclude a previously found suboptimal solution. CDSAT is designed for SMA with SMT as a special case [42, 7]. CDSAT is defined as a *transition system* that orchestrates in a conflict-driven manner *theory inference systems*, called *theory modules*. A theory module is an abstraction of a theory reasoning procedure. The candidate model is represented by a *trail* Γ of assignments, which includes the input and is shared by all theory modules. Assignments in Γ are either *decisions* or *justified assignments*, where the *justification* is a set of prior assignments in Γ . Decisions can be either Boolean or first-order. Input assignments are justified assignments with empty justification. All justified assignments are Boolean except for the input first-order assignments of an SMA problem [42, 7].

The CDSAT Transition System. The *Decide* rule allows a theory module to propose an *acceptable* assignment for a term that is *relevant* for its theory. The *Deduce* rule adds to the trail Γ a justified assignment based on a theory inference, provided the term of the derived assignment

comes from a *finite global basis*. Deductions cover both *propagations* and inferences that *detect* and *explain* theory conflicts, letting them surface in Γ as Boolean conflicts. If the conflict is at level 0, rule *Fail* reports unsatisfiability. Otherwise, rule *ConflictSolve* passes control to the *conflict state rules*, returning the trail they produce. A *conflict state* is given by a trail Γ and a conflict, which is an unsatisfiable subset of Γ , containing in general both Boolean and first-order assignments. The *Resolve* rule unfolds the conflict, replacing a justified assignment by its justification, until the *Backjump* rule can solve the conflict by flipping a Boolean assignment, so that the procedure will not hit the same conflict. In the Boolean case, these two rules can emulate conflict solving as in CDCL. First-order assignments cannot be flipped. *UndoClear* solves the conflict by undoing a first-order decision A and clearing Γ of its consequences, when Γ contains a *late propagation* that makes A unacceptable, so that it will not be repeated. Otherwise, the *UndoDecide* rule solves the conflict by undoing a first-order decision A , clearing Γ of its consequences, and flipping a Boolean one, so that A will not be retried. If the theories are *disjoint*, there is a finite global basis that contains the input, and the theory modules are sound and *leading-theory complete*, CDSAT is *sound*, *terminating*, and *complete* [42, 7]. These properties are preserved if the *Backjump* rule is replaced by a more general *LearnBackjump* rule, that allows to flip a Boolean subset of the conflict into a learned clause [40, 6]. We also extended CDSAT with *proof generation* towards different proof formats, including resolution-based proofs [40, 6].

CDSAT Theory Modules. The inference rules of a theory module derive Boolean assignments from assignments, and can generate *new* (i.e., non-input) terms, provided they come from a finite *local basis* for that theory. We defined theory modules and local bases for *propositional logic*, and for the quantifier-free fragments of the theories of *equality*, *linear rational arithmetic*, and *arrays with extensionality* [42, 7, 6]. A theory procedure that is not conflict-driven is integrated in CDSAT by viewing it as a *black-box theory module*, whose only inference rule invokes the procedure to determine that a set of literals is unsatisfiable [42, 7, 6]. If all modules are black-boxes, CDSAT can emulate equality sharing [6]. However, CDSAT does not require stable infiniteness, provided there is a *leading theory* \mathcal{T}_1 that knows all sorts in the union of theories and acts as an aggregator of cardinality requirements by different theories. The theory module of the leading theory enforces the aggregated requirements, such as *at-most cardinality constraints* [6]. For all above mentioned theories \mathcal{T} , we proved that the \mathcal{T} -module is *leading-theory complete* [6]. This means that if the \mathcal{T} -module cannot expand an assignment, for all \mathcal{T}_1 -models satisfying the assignment there is a satisfying \mathcal{T} -model that agrees with the \mathcal{T}_1 -model on cardinality of shared sorts and equality of shared terms [7]. We also showed how a finite *global basis* can be built from the local bases [6].

CDSAT for Nondisjoint Theories with Shared Predicates. We are extending the CDSAT framework to *predicate-sharing unions*, that is, unions of theories that are *either disjoint or share only predicate symbols* (in addition to equality) [35, 1]. Consider a *theory of arrays with length*, where extensionality says that two arrays are equal if they have the same length n and the same elements at all indices between 0 and $n - 1$. In order to write this axiom, one needs symbols from linear integer arithmetic (LIA), so that the two theories are nondisjoint. Also, such an axiomatization forces the indices to be integers, which is not imposed by the theory of arrays

without length. We proposed a *theory of arrays with abstract length*, where the notion of an index being within bounds is abstracted into that of an index being *admissible* [35]. The admissibility predicate is shared by the theory of arrays and another theory, which may be LIA, but does not have to. The admissibility predicate is free in the theory of arrays and interpreted in the other. This approach covers several interpretations of length and admissibility, including one where length is given by starting address in memory and number of admissible indices. The admissibility predicate is the only symbol that the theories need to share. The only definitions of the CDSAT framework that need to be generalized to accommodate shared predicates are those of *relevance* (of a term to a theory for the purpose of decisions) and *leading-theory completeness*. We gave a theory module for the theory of arrays with abstract length, and we proved that it is leading-theory complete [35]. While soundness and termination are unaffected by the generalization, we proved that CDSAT is *complete for predicate-sharing unions*. We are working on modules for the theories of *maps* (e.g., *hashmaps*) and *dynamic arrays* (aka *vectors*) with abstract length, and on generalizing the global basis construction to predicate-sharing unions [1].

B. Interpolation of Proofs

Interpolation is an automated reasoning technique that finds application in abstraction refinement, safety checking, and invariant generation. Given two disjoint sets of clauses A and B , such that $A \cup B$ is inconsistent, a (reverse) *interpolant* of (A, B) is a formula that is implied by A , inconsistent with B , and such that its uninterpreted symbols are common to A and B . If B encodes a partial model, a reverse interpolant is a candidate *explanation* of why A is in conflict with B . Therefore, interpolation is relevant to conflict-driven satisfiability procedures [39].

B1. Interpolation of Ground Proofs by Superposition. A complete *interpolation system* for an inference system Γ extracts an interpolant of (A, B) from any Γ -refutation of $A \cup B$. It works by attaching a *partial interpolant* to every clause in the refutation, in such a way that the partial interpolant of \square is an interpolant of (A, B) . For each inference rule the partial interpolant of the conclusion is defined *inductively* from those of the premises. In a proof by propositional resolution (hence CDCL), all literals are input literals, hence either *A-colored* (the symbol occurs in A but not in B), *B-colored* (the symbol occurs in B but not in A), or *transparent* (the symbol occurs in both). Thus, the partial interpolant of the resolvent is defined based on whether the literal resolved upon is *A-colored*, *B-colored*, or *transparent* [45, 10]. In a first-order proof with equality, even in the ground case, new literals are generated. Assume that an *AB-mixed* equality $t_a \simeq t_b$ is generated, where terms t_a and t_b are in normal form, t_a is *A-colored* (made of *A-colored* and *transparent* symbols), and t_b is *B-colored* (made of *B-colored* and *transparent* symbols). If $t_a \succ t_b$, all occurrences of t_a should be rewritten to t_b , or the congruence classes of t_a and t_b should be merged with t_b as representative, jeopardizing a case analysis based on colors. I gave a superposition-based proof that the quantifier-free fragment of the theory of equality is *equality-interpolating* (if $t_a \simeq t_b$ holds, then also $t_a \simeq t \wedge t_b \simeq t$ for some transparent t holds). Also, I showed that an ordering where transparent terms are smaller than the others guarantees that ground superposition proofs do not contain *AB-mixed* literals [10]. Then, we designed *the first complete interpolation system* for *ground* refutations by superposition [79, 10].

B2. Interpolation of Non-Ground Proofs by Superposition and CDCL($\Gamma + \mathcal{T}$). In non-ground proofs, AB -mixed literals are unavoidable, even when the only colored symbols are constants, because matching substitutions and most general unifiers mix the symbols. Therefore, we designed a *two-stage approach* [89, 11]. In the first stage, a *provisional interpolation system* computes a *provisional interpolant*, that is entailed by A and inconsistent with B , but may contain non-shared symbols. We defined a *complete provisional interpolation system*, for an ordering-based inference system Γ with resolution and superposition, that produces provisional interpolants where all predicate symbols are transparent [11]. In the second stage, colored constants are replaced with quantified variables (*lifting*). I proved that the lifting of a provisional interpolant is an interpolant, so that the two-stage approach yields *the first complete interpolation system* for non-ground Γ -refutations, provided the only colored symbols in the provisional interpolant are constants [11]. By combining the provisional interpolation system for Γ with one for CDCL(\mathcal{T}), we get a provisional interpolation system for CDCL($\Gamma + \mathcal{T}$), so that lifting yields interpolants for CDCL($\Gamma + \mathcal{T}$)-refutations. The two-stage approach can interpolate refutations by equality sharing, CDCL(\mathcal{T}), and CDCL($\Gamma + \mathcal{T}$), even if there is a theory that is not convex or not equality-interpolating. It also handles the *model-based theory combination* variant of equality sharing (used in CDCL($\Gamma + \mathcal{T}$)), where a model-constructing \mathcal{T} -satisfiability procedure propagates equalities that are true in the candidate \mathcal{T} -model rather than entailed.

C. Distributed Automated Deduction

I was the first one to investigate *distributed automated deduction* [108, 66, 64, 63, 28, 62, 29, 60, 61, 25, 27, 97, 22, 23, 57, 56, 84, 83, 18, 54, 73]. I analyzed the parallelizability of theorem-proving strategies, classifying types of parallelism based on the granularity of data accessed in parallel: *fine-grain parallelism* is *parallelism at the term/literal level*, *medium-grain parallelism* is *parallelism at the clause level*, and *coarse-grain parallelism* is *parallelism at the search level* [108, 29, 83, 18, 73]. Parallelism at the term/literal level affects operations *below the inference or clause level* (e.g., *parallel rewriting*), but it requires clause preprocessing, which is problematic in theorem proving, where new clauses are generated. Parallelism at the clause level yields *parallel inferences*, but it is at odd with *eager contraction*, as priority to contraction reduces the concurrency of inferences. The possibility of *conflicts* between parallel inferences is an obstacle especially for *contraction-based* strategies, where *backward contraction*, the contraction of pre-existing clauses by new ones, applies to clauses active as premises of expansion inferences. Thus, I proposed *parallelism at the search level* by *Clause-Diffusion* [108, 66, 64, 63, 28, 62, 29, 25, 27].

C1. The Clause-Diffusion Method. In Clause-Diffusion, multiple deductive processes *search in parallel* the space of the problem and cooperate to seek a proof [108, 25, 27, 60, 97, 22, 73]. All processes start with the same input problem, ordering-based inference system, and search plan, although different search plans may be assigned. Every process develops *its own derivation* and builds *its own database of clauses* independently. The processes are *asynchronous*, as the only synchronization occurs when one sends all others a halting message because it found a proof. Clause-Diffusion is a *distributed-search* method, because it *subdivides the search space* by subdividing clauses and inferences. In an ordering-based strategy, a newly generated clause φ

is subject to *forward contraction*, the contraction of new clauses by pre-existing ones. If the resulting normal form $\varphi \downarrow$ is not trivial, it is kept. In Clause-Diffusion, whenever a process generates and keeps a clause, it assigns it to a process, possibly to itself, by an *allocation criterion*. Every clause is *owned* by a process, and since every clause has its own variables, and variants are distinct clauses, every clause is owned by *only one* process. I designed several *heuristic allocation criteria* for this purpose [56, 73]. Then, expansion inferences are subdivided based on ownership of the premises: for example, a process paramodulates only into the clauses it owns. Backward contraction inferences that generate clauses are subdivided without delaying the deletion of redundant clauses: whenever a process detects that a clause φ can be backward-simplified, it deletes it, but generates $\varphi \downarrow$ only if it owns φ . Every kept clause, regardless of whether generated by either expansion or backward contraction, is given a *unique global identifier* and is broadcast as an *inference message* for completeness, hence the name of the method.

C2. Properties of the Clause-Diffusion Method. I defined *fairness of distributed derivations*, giving sufficient conditions and showing that Clause-Diffusion satisfies them [108, 64, 25]. Thus, if the inference system is refutationally complete and the search plan at each process is fair, parallelization by Clause-Diffusion preserves *completeness*. I discovered that *subsumption in distributed derivations* may violate fairness and the soundness of contraction; I provided a general solution that preserves these properties without renouncing subsumption [108, 28]. Clause-Diffusion also achieves *distributed global contraction* and *distributed proof reconstruction* [22, 73]. The former property ensures that if φ is globally redundant at some stage of the distributed derivation, φ is recognized redundant eventually by every process. The latter property ensures that the process that generates \square is able to reconstruct the proof from the final state of its database, even if all processes contributed to the proof. I gave sufficient conditions for this property, and proved that Clause-Diffusion fulfills them, without centralized control, or ad hoc postprocessing [22].

C3. The Clause-Diffusion Provers and Super-Linear Speed-Up. I implemented several *Clause-Diffusion theorem provers*. *Aquarius* parallelized Bill McCune’s OTTER 2.2 theorem prover for FOL [66, 108, 63, 27]. *Peers* was built on top of a prototype prover (from Bill’s Otter Parts Store) for equational theories modulo associativity and commutativity (AC) [62, 25]. *Peers-mcd* implemented *Modified Clause-Diffusion* [22], the final version of the methodology. *Peers-mcd.a* [22] had the same sequential basis as *Peers*. All subsequent versions parallelized Bill’s EQP prover for equational theories modulo AC. EQP became famous in 1996 for proving that *Robbins algebras are Boolean*, a conjecture open since 1933. In most experiments, at least one allocation criterion allowed *Peers-mcd.b* to speed-up over EQP’s best performance, with *super-linear speed-up* in two thirds of the proof of the Robbins theorem [57, 56]. Clause-Diffusion enables super-linear speed-up, because it does not compute in parallel the sequential search, but it uses distributed search to generate different searches. *Peers-mcd.b* also generated the *first mechanical proof* of the *Levi commutator problem* [84]. The proof of the Robbins theorem by *Peers-mcd.c* was *the fastest* at the time [18]. *Peers-mcd.d* [54] featured both *distributed search* and *multi-search*, where the processes apply different search plans, including *target-oriented heuristics* [106, 69, 65]. *Peers-mcd.d* offered *distributed-search* strategies (search space subdivided and same search plan for all processes),

multi-search strategies (no subdivision and different search plans), and *hybrid* ones (subdivision and different search plans). I investigated whether Peers-mcd.d could prove the *Moufang identities* without building cancellation laws in the inference system. With some strategies, EQP could not find a proof while Peers-mcd.d did. With others, EQP succeeded, but Peers-mcd.d was faster, with instances of *super-linear speed-up*. Distributed search behaved better than multi-search, which did not find the proofs, and their hybridization performed even better [54].

C4. PSATO. Clause-Diffusion inspired *PSATO*, the first distributed-search SAT-solver with a *divide-and-conquer* organization [61, 23, 73]. A master process *partitions* the search space, by assigning *disjoint subproblems* to slave processes, each executing the Davis-Putnam-Logemann-Loveland procedure for SAT. Each subproblem is defined by a *guiding path*, which encodes a Boolean assignment. Thus, every subproblem is a Boolean instance of SMA. A guiding path can be read as a conjunction of literals, later called a *cube*, so that PSATO is an ancestor of the *cube-and-conquer* approach to parallel CDCL-based SAT solving. PSATO solved *quasigroup existence problems*, including some that were never conquered before [61, 23].

D. Strategy Analysis

Theorem-proving strategies are evaluated by comparing the performances of their implementations. Worst-case or average-case analyses do not apply, as theorem-proving strategies are only semidecision procedures. The search space is infinite, and the complexity of searching for a proof is proportional to neither input nor output size [96, 95, 105]. If an empirical evaluation indicates that a feature is useful, it remains the question of why. An intuitive explanation may say “contraction helps by pruning the search space,” but deleting finitely many branches in an infinite search graph does not make it finite. How do we compare infinite spaces to say that one is “smaller”? Which computational complexity is affected, when the derivation may not halt, so that “time” is not defined? Such questions led me to think about formal tools for *strategy analysis*.

D1. Analysis of Ordering-Based Strategies. I introduced the *marked search-graph* as a *model of the search space and search process* that covers both expansion and contraction inferences [94, 58, 93, 20]. In this model, the *search graph* represents the space of all possible inferences, and the *marking* represents the *search process*, with generations and deletions of clauses. A *complexity measure* involves a *well-founded ordering* and representative objects to be compared. I observed that at each stage of a derivation a finite portion of the search space has been generated (the *present*) and an infinite portion remains to be explored (the *future*). While a strategy works with a finite amount of data, capturing the complexity of a search problem requires to measure changes *in both present and future*, since that finite amount of data can generate anything in the future. Since the latter is infinite, one needs to impose a *bound*. I used the *ancestor-graph* of a clause to define its *distance* from the input. The *bounded search space* with bound j is the *multiset of clauses* reachable within distance j , where the multiplicity of a clause is the number of ancestor-graphs of its variants within distance j . The infinite search space is treated as an *infinite succession* (for all j) of bounded search spaces. Since they are finite, the bounded search spaces

can be compared by the multiset extension of a well-founded ordering on clauses or proofs (the ancestor-graphs). I analyzed *contraction-based strategies of different contraction power*, showing that more contraction eventually causes a bigger reduction of the bounded search spaces [58, 20].

D2. Analysis of Distributed Ordering-Based Strategies. In *distributed search* multiple processes are active in parallel. I devised the *parallel marked search-graph* to model also the *subdivision* of the search space among the processes, the effects of *communication*, and the *overlap* of the processes [85, 92, 55, 19]. The bounded search spaces are defined relative to each process, with the multiplicity of a clause given by the number of ancestor-graphs within distance j allowed to that process by the subdivision scheme. In the *parallel bounded search spaces*, the multiplicity of a clause is the *average* of its multiplicities at the processes, so that the overlap is taken into account. Subdivision and contraction make the bounded search spaces smaller, whereas communication undoes in part this impact. I compared a distributed-search contraction-based strategy with its sequential basis, and analyzed the overlaps due to inaccurate subdivision and to communication, giving *sufficient conditions* to avoid the first and minimize the second. Then, I discovered two patterns of worst-case behavior, called *late contraction* and *contraction undone*, where the interaction of asynchronous communication and contraction violates eager contraction. It follows that sufficient conditions for the parallel bounded search spaces to be smaller than or equal to the sequential ones are minimum overlap and immediate propagation of clauses. Counterexamples show that weaker assumptions are not sufficient. Since these conditions are *not necessary*, distributed-search contraction-based strategies may still behave well in practice, and even exhibit *super-linear speed-up's* [56], approximating the ideal behavior in the theorem.

D3. Analysis of Subgoal-Reduction Strategies. While the search space of an ordering-based strategy is described by a *synthetic search-graph*, where vertices are labelled by clauses and arcs capture the *synthesis* of a new clause from existing ones, the search space of a tableau-based strategy is described by an *analytic search-graph*, where vertices are labelled by literals and arcs capture the *decomposition* of clauses into literals. I defined *synthetic marked search-graphs* for linear resolution and *analytic marked search-graphs* for *clausal normal form tableaux*, including *model-elimination tableaux*. In the analytic marked search-graph the marking captures the application of substitutions to *rigid variables*, hence to the whole tableau, the closure of branches, and the effects of backtracking [17]. The *distance* of a vertex from the root is the length of its *ancestor-path*. Since an ancestor-path is labeled by a *sequence of literals*, which represents a *partial interpretation*, the *bounded search spaces* are *multisets of partial interpretations*. The multiplicity of an interpretation is the number of ancestor-paths labelled by that interpretation within the bound on distance. Because of the bound, these multisets are finite, and can be compared by comparing their *cardinalities*, which is suitable for strategies that survey and eliminate candidate models. I analyzed tableau-based strategies with and without the *regularity check*, that prevents the repetition of literals on a branch, and with and without *lemmaizing by folding-up*, showing that both refinements reduce the bounded search spaces [17].